

# Basic Deep Learning in Computer Vision

---

## **Day 2**

Morning: Introduction to Artificial Neural Network

Afternoon: Introduction to Convolution Neural Network



# Programme

---

	Morning	Afternoon
<b>Day 1</b>	<ul style="list-style-type: none"><li>• Computer Vision</li><li>• Image Libraries<ul style="list-style-type: none"><li>• Activity 1: Getting Started with Libraries</li></ul></li></ul>	<ul style="list-style-type: none"><li>• Image Preprocessing</li><li>• Image Augmentation<ul style="list-style-type: none"><li>• Activity 2: Image preprocessing</li><li>• Activity 3: Image Augmentation</li></ul></li></ul>
<b>Day 2</b>	<ul style="list-style-type: none"><li>• Basic of Neural Network<ul style="list-style-type: none"><li>• Activity 4: Building NN with Python</li></ul></li><li>• Introduction to Keras<ul style="list-style-type: none"><li>• Activity 5: Building NN with Keras</li></ul></li></ul>	<ul style="list-style-type: none"><li>• Image Convolution</li><li>• Convolution Neural Network (CNN)<ul style="list-style-type: none"><li>• Activity 6: Create and use CNN</li></ul></li><li>• Quiz</li></ul>

# Artificial Neural Network (ANN)

---



# Definition of ANN

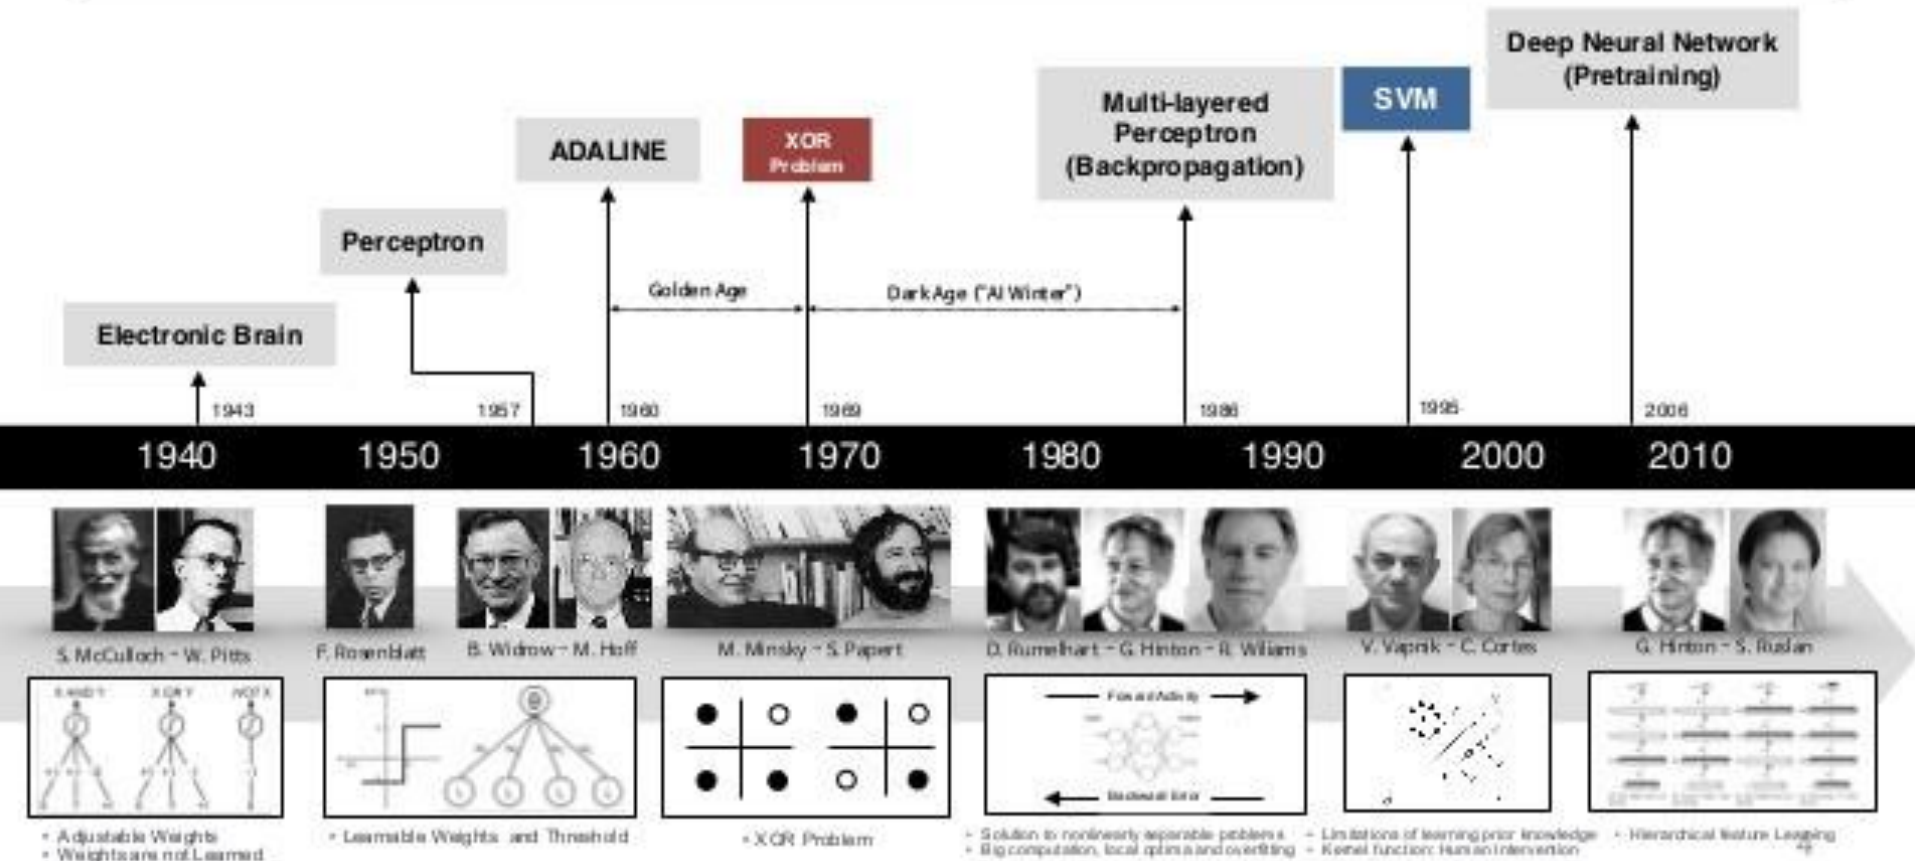
- An ANN is based **on a collection of connected units or nodes called artificial neurons**, which loosely model the neurons in a biological brain. Each connection, like the synapses in a biological brain, can **transmit a signal to other neurons**. An artificial neuron that **receives a signal then processes it and can signal neurons connected to it**. The "signal" at a connection is a real number, and the output of each neuron is computed by some non-linear function of the sum of its inputs.
- The connections are called edges. Neurons and edges typically have a **weight that adjusts as learning proceeds**. The weight increases or decreases the strength of the signal at a connection. **Neurons may have a threshold** such that a signal is sent only if the aggregate signal crosses that threshold.
- Typically, neurons are aggregated into **layers**. Different layers may perform different transformations on their inputs. Signals travel from the first layer (the **input layer**), to the last layer (the **output layer**), possibly after traversing the layers multiple times.
- Ref: [https://en.wikipedia.org/wiki/Artificial\\_neural\\_network](https://en.wikipedia.org/wiki/Artificial_neural_network)



# Brief History of ANN

## Brief History of Neural Network

DEVIEW  
2015





Ref: <https://medium.com/coinmonks/neural-networks-bb11fb9a8266>



# Inspiration from neural science

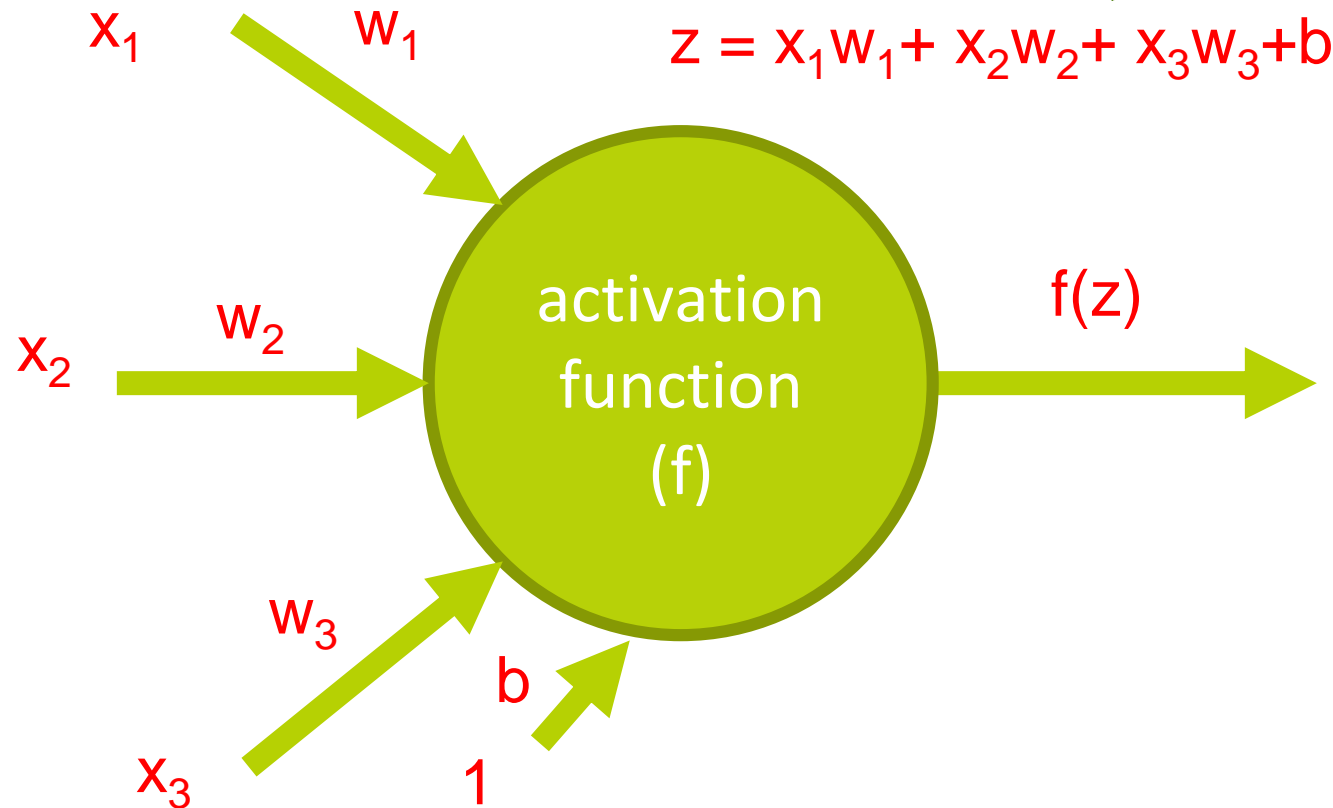
- A **neuron** takes in several signals through **dendrites**, processes the signals and sends out a signal (spike of electrical activity) through the **axon**.
- The axon is split into several branches and connects to other dendrites through **synapse**, forming the biological neural network.





# A Neural Node

This is similar to  
the formula of  
Linear Regression





# Activation function

---

- Taking inspiration from the biological neuron, there is a need to '***threshold***' the output signal. Such function is known as ***Activation Function***.
- Most commonly used function includes **Sigmoid, Tanh** and **Relu**






# Activation function

- **Sigmoid ( $\sigma$ )**

- Limit to 0 to 1


$$f(z) = \frac{1}{1 + e^{-z}}$$



- **Tanh – Hyperbolic tangent**

- Limit to -1 and 1

$$\tanh(z) = \frac{\sinh(z)}{\cosh(z)} = \frac{e^{2x} - 1}{e^{2x} + 1}$$




- **Relu – Rectified Linear Unit**

- linear positive values
- Ignore negative value

$$ReLU(z) = \begin{cases} 0, & z < 0 \\ z, & z \geq 0 \end{cases}$$

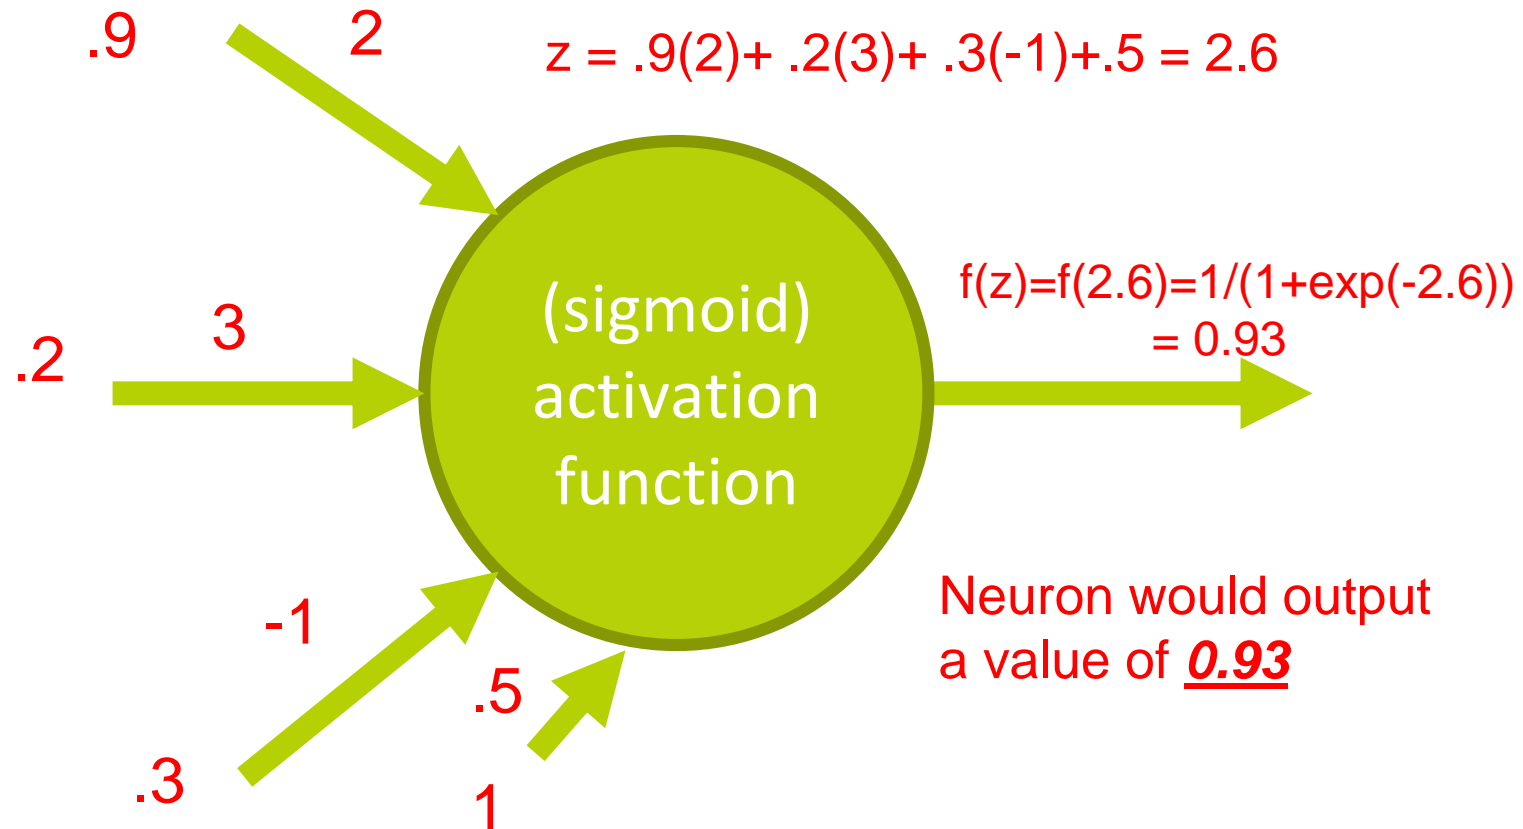
t, ACE@RP





# An example of a node


- When Sigmoid is used, the output equation of a neuron is similar to the *Logistic Regression*. Hence, a neuron **node** is also interpreted as a **classifier**.





# Forming a hidden layer


- Placing several nodes side by side, a layer is formed





# Forming a hidden layer

- Putting the equations together as a matrix





# Input and Output Layer


---

- Input layer allows the features (or data) to enter the neural network.
  - The number of input nodes must be equal to the number of features
  - The output of the layer is equal to the feature values
- Output layer produces the predicted labels.
  - The number of output nodes must be equal to the number of labels
  - The output of the output layer is calculated similarly to the hidden layer
  - It has an additional activation function called **Softmax**.



# Forming a network


- One input layer  $\rightarrow$  zero to many hidden layer(s)  $\rightarrow$  one output layer
- The bias is normally not drawn or shown in a network diagram





# Feedforward

- Defining the values






# Feedforward

- Calculate the output

$$\begin{aligned} z^{(2)} &= a^{(1)}W^{(1)} & z^{(3)} &= a^{(2)}W^{(2)} \\ a^{(2)} &= \sigma(z^{(2)}) & a^{(3)} &= \sigma(z^{(3)}) \end{aligned}$$

$$z^{(4)} = a^{(3)}W^{(3)}$$

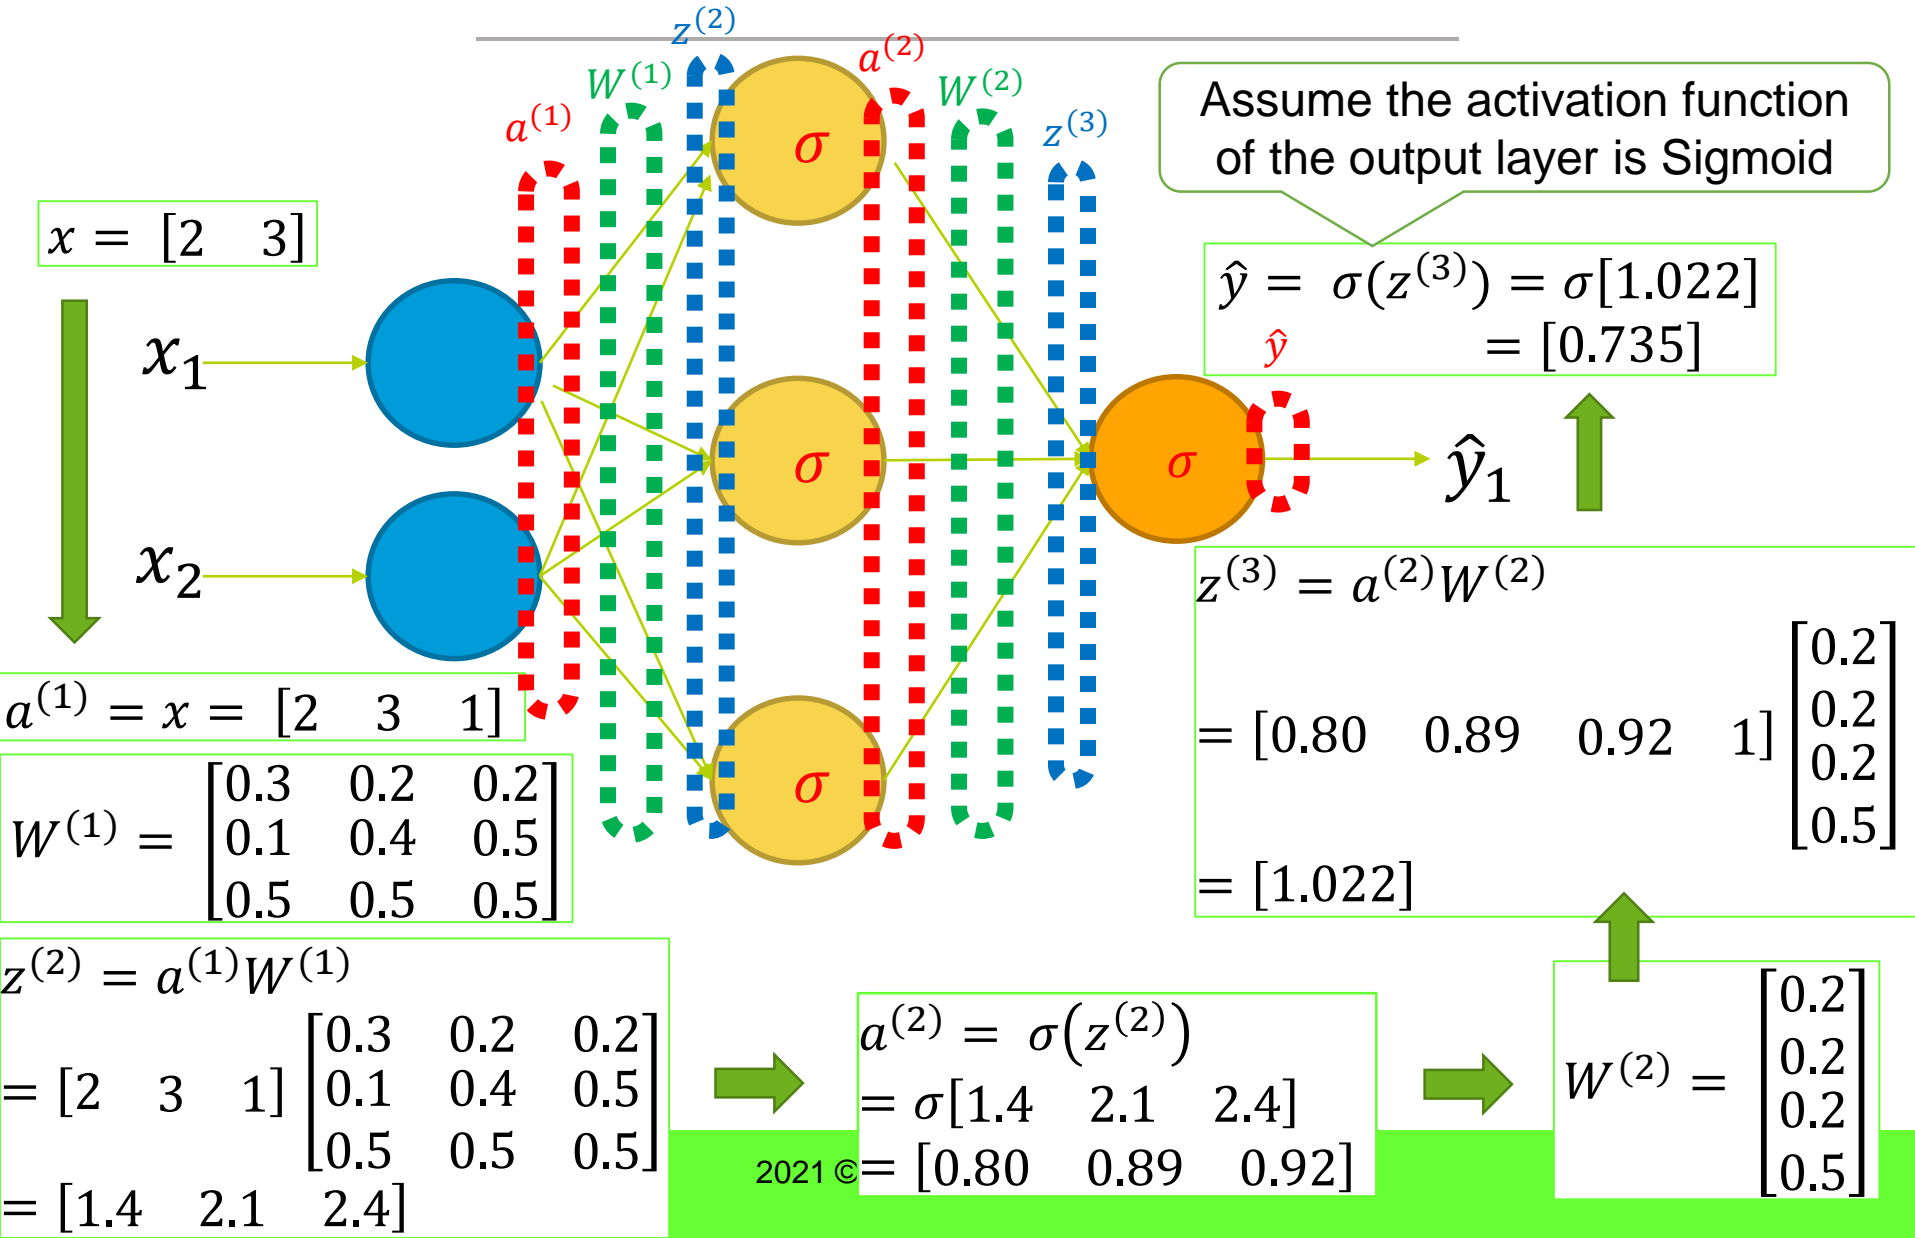
$$\hat{y} = \textit{softmax}(z^{(4)})$$







# Feedforward example





# Feedforward example

- Multiple inputs can be calculated in just **one** feedforward pass. The number of inputs is known as **Batch Size**.

$$\begin{aligned}
 x = \begin{bmatrix} 2 & 3 \\ 2 & 4 \\ \vdots & \vdots \\ 1 & 4 \\ 9 & 4 \end{bmatrix} & \xrightarrow{\quad} \dots = \begin{bmatrix} 2 & 3 & 1 \\ 2 & 4 & 1 \\ \vdots & \vdots & \vdots \\ 1 & 4 & 1 \\ 9 & 4 & 1 \end{bmatrix} \begin{bmatrix} 0.3 & 0.2 & 0.2 \\ 0.1 & 0.4 & 0.5 \\ 0.5 & 0.5 & 0.5 \end{bmatrix} \dots \xrightarrow{\quad} \hat{y} = \begin{bmatrix} 0.735 \\ 0.738 \\ \vdots \\ 0.736 \\ 0.748 \end{bmatrix} \\
 & = \begin{bmatrix} 1.4 & 2.1 & 2.4 \\ 1.5 & 2.5 & 2.9 \\ \vdots & \vdots & \vdots \\ 1.2 & 2.3 & 2.7 \\ 3.6 & 3.9 & 4.3 \end{bmatrix}
 \end{aligned}$$

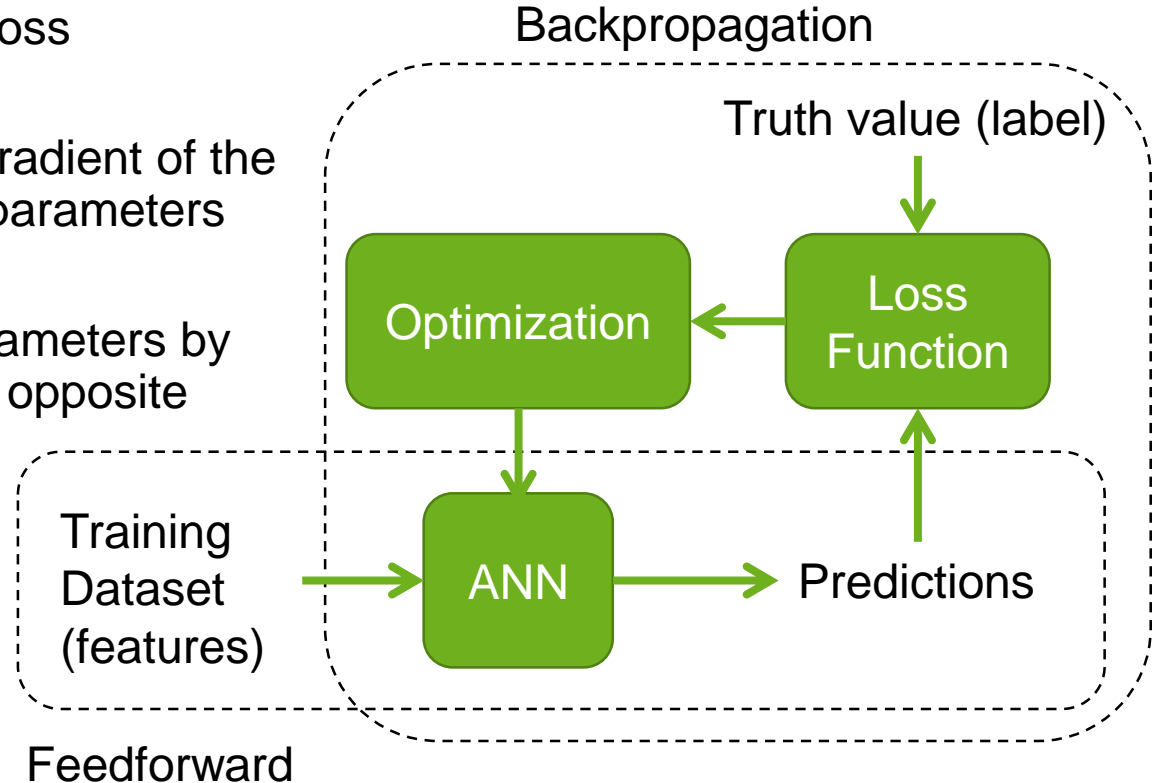
The diagram illustrates a feedforward neural network calculation. It shows an input vector  $x$  (a 5x2 matrix) being multiplied by a weight matrix  $W^{(1)}$  (a 5x3 matrix) to produce a hidden layer output  $z^{(2)}$  (a 5x3 matrix). The hidden layer output is then passed through an activation function  $a^{(1)}$  to produce the final output  $\hat{y}$  (a 5x1 vector). The activation function is applied element-wise to the hidden layer output, as shown by the second matrix in the sequence.



# Training of weights

- **Work flow**

- Step 1: Make a prediction (feedforward)
- Step 2: Calculate Loss
- Step 3: Calculate gradient of the loss function w.r.t. parameters
- Step 4: Update parameters by taking a step in the opposite direction
- Step 5: Iterate





# Training of weights

- **Loss Function,  $J(y, \hat{y})$** , calculate how **poorly** our model is performing by comparing what the **model is predicting** with the **actual value** it is supposed to output.

- Regression Loss – the model is predicting a continuous value

- Mean Square Error
- Mean Absolute Error

$$Loss = \frac{1}{n} * \sum_{i=0}^n (Y_i - Y_{pred_i})^2$$

- Classification Loss – the model is predicting a discrete value or specific class.
  - Binary Cross Entropy
  - Categorical Cross Entropy

$$Loss = (Y)(-\log(Y_{pred})) + (1 - Y)(-\log(1 - Y_{pred}))$$

Ref: <https://deeplearningdemystified.com/article/fdl-3>



# Training of weights

- **Optimization** - Optimizers are algorithms or methods used to change the attributes of your neural network such as weights and learning rate in order to reduce the losses.

- Stochastic Gradient Descent (SGD)

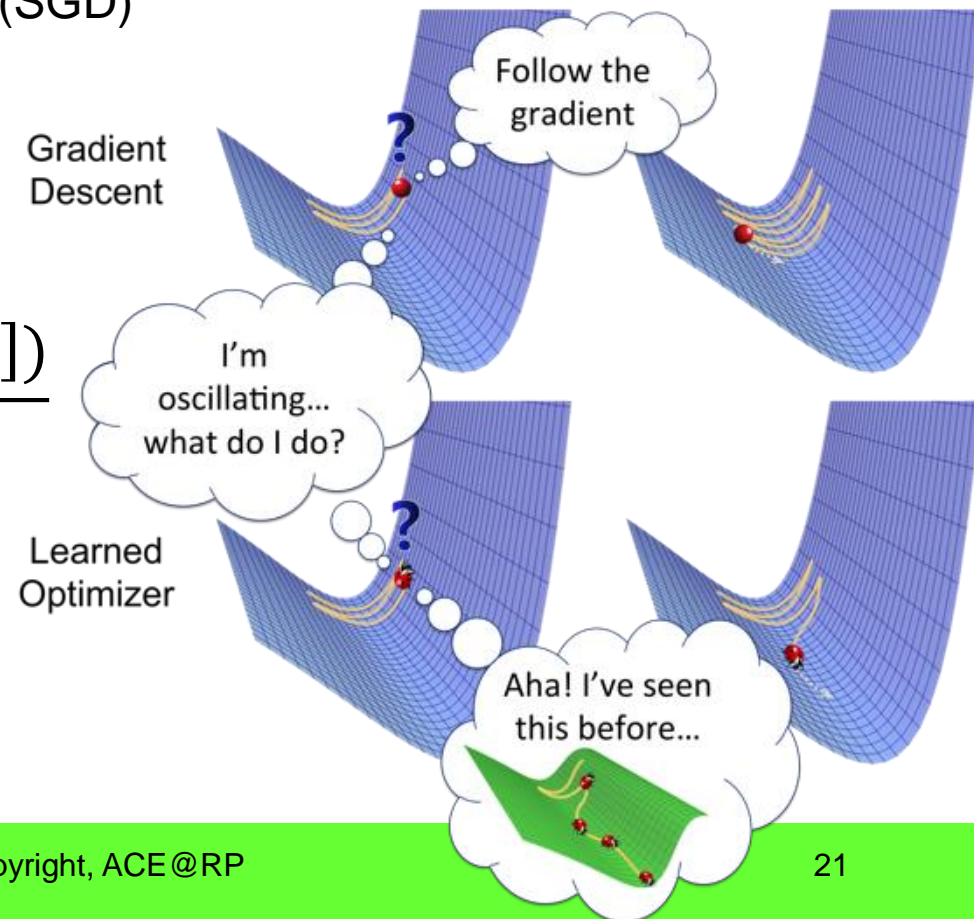
$$w = w - \alpha \frac{\partial J(y, \hat{y})}{\partial w}$$

Gradient Descent

- Mini-Batch Gradient Descent

$$w = w - \alpha \frac{\partial J([y], [\hat{y}])}{\partial w}$$


- Momentum
- Adam
- RMSProp



Ref: <https://towardsdatascience.com/optimizers-for-training-neural-network-59450d71caf6>



# Training of weights (An Example)



Current Values:

$$x = 2, y = 4$$

$$w = 0.2, b = 0.5$$

- **Assume:**

- Activation function = Linear
- Loss function = MSE
- Optimization = SGD (one sample used)
- Learning rate = 0.1

- **Step 1: Make a prediction**

$$\begin{aligned}\hat{y} &= xw + b \\ &= 2 * 0.2 + 0.5 \\ &= 0.9\end{aligned}$$



# Training of weights (An Example)

---

- **Step 2: Calculate the Loss**

$$\begin{aligned} J(y, \hat{y}) &= (y - \hat{y})^2 \\ &= (4 - 0.9)^2 \\ &= 9.61 \end{aligned}$$

$$\begin{aligned} J(y, \hat{y}) &= (y - \hat{y})^2 \\ &= y^2 - 2y\hat{y} + \hat{y}^2 \\ &= y^2 - 2y(xw - b) + (xw - b)^2 \\ &= y^2 - 2xyw + 2by + x^2w^2 - 2bxw + b^2 \\ &= x^2w^2 - (2bx + 2xy)w + (y^2 + 2by + b^2) \\ &\text{or} \\ &= b^2 + (2y - 2xw)b + (y^2 - 2xyw + x^2w^2) \end{aligned}$$




# Training of weights (An Example)

- Step 3: Calculate the gradient**

$$\begin{aligned}
 J(y, \hat{y}) &= (y - \hat{y})^2 \\
 &= y^2 - 2y\hat{y} + \hat{y}^2 \\
 &= y^2 - 2y(xw - b) + (xw - b)^2 \\
 &= x^2w^2 - (2bx + 2xy)w + (y^2 + 2by + b^2)
 \end{aligned}$$

$$\begin{aligned}
 \frac{\partial J}{\partial w} &= 2x^2w - 2x(b + y) \\
 &= (2)(2 * 2)(0.2) - (2)(2)(0.5 + 4) \\
 &= 1.6 - 18 \\
 &= -16.4
 \end{aligned}$$

$$J = c_1 w^2 + c_2 w + c_3$$







# Training of weights (An Example)

---

- **Step 4: Update the weights**

$$\begin{aligned}w &= w - \alpha \frac{\partial J(y, \hat{y})}{\partial w} \\&= 0.2 - (0.1)(-16.4) \\&= 1.84\end{aligned}$$

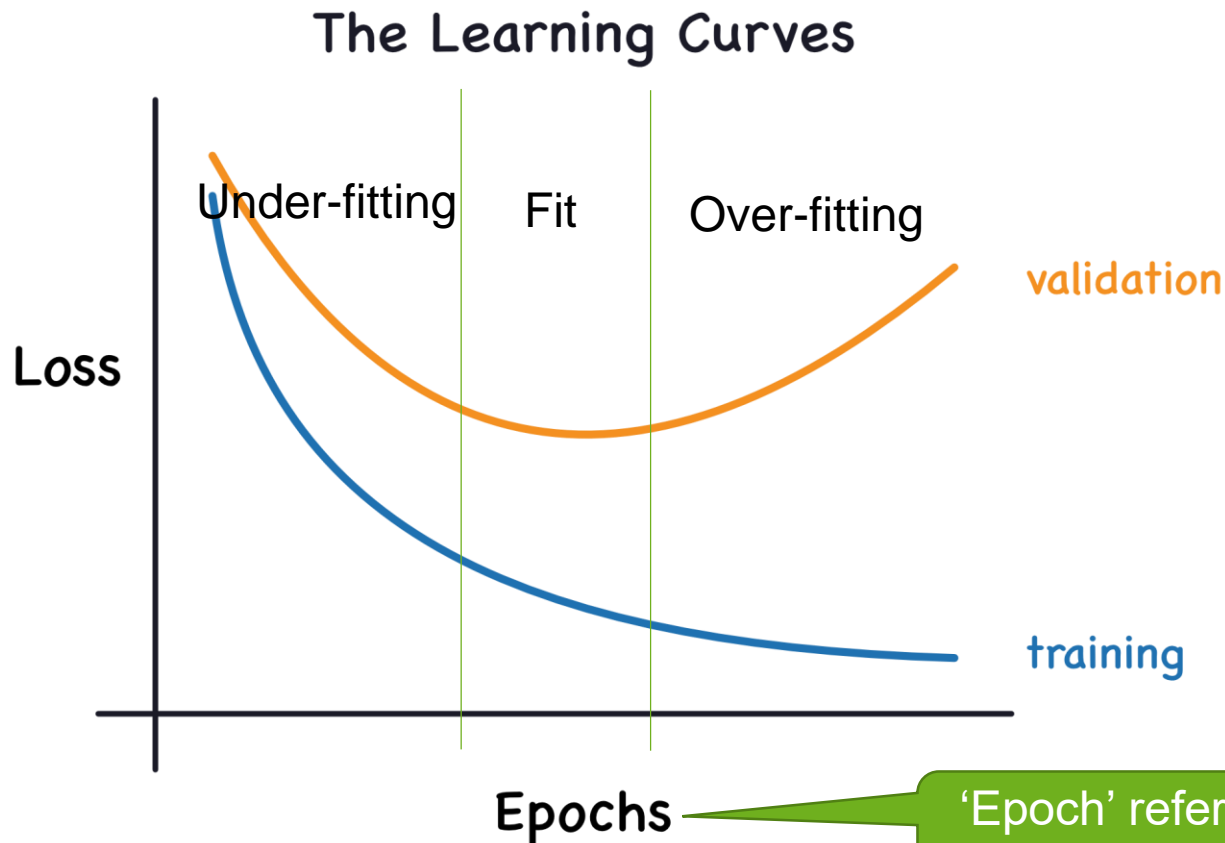
Try out the same approach  
to update the bias.



# Training of weights (An Example)

- **Step 5: Iterate**

- A learning curve will be produced




Ref:  
<https://www.kaggle.com/ryanholbrook/overfitting-and-underfitting>



# Backpropagation

- To train the weights of a neural network, we need to find the gradient of each weights.
- This is very difficult.

Want:  $\frac{\partial J(y_i, \hat{y}_i)}{\partial W_k}$





# Backpropagation

- Using **Chain Rule** and **Calculus**
- The calculated values can be “propagate back” from the higher layers to the lower layers.
- Although the formula is long, but it can be quickly calculated by the computer.

$$\frac{\partial J(y, \hat{y})}{\partial W^{(3)}} = \underbrace{(\hat{y} - y)} \cdot a^{(3)}$$

$$\frac{\partial J(y, \hat{y})}{\partial W^{(2)}} = \underbrace{\left[ (\hat{y} - y) \cdot W^{(3)} \cdot \sigma'(z^{(3)}) \right]} \cdot a^{(2)}$$

$$\frac{\partial J(y, \hat{y})}{\partial W^{(1)}} = \underbrace{\left[ (\hat{y} - y) \cdot W^{(3)} \cdot \sigma'(z^{(3)}) \right] \cdot W^{(2)} \cdot \sigma'(z^{(2)})}_{\cdot X}$$


Where:  $\sigma'(z) = \sigma(z)(1 - \sigma(z))$



# Backpropagation (optional)

---

- **Explain by YouTube channel - 3Blue1Brown**
  - Check out the series on Neural Network.
  - [https://www.youtube.com/playlist?list=PLZHQObOWTQDNU6R1\\_67000Dx\\_ZCJB-3pi](https://www.youtube.com/playlist?list=PLZHQObOWTQDNU6R1_67000Dx_ZCJB-3pi)



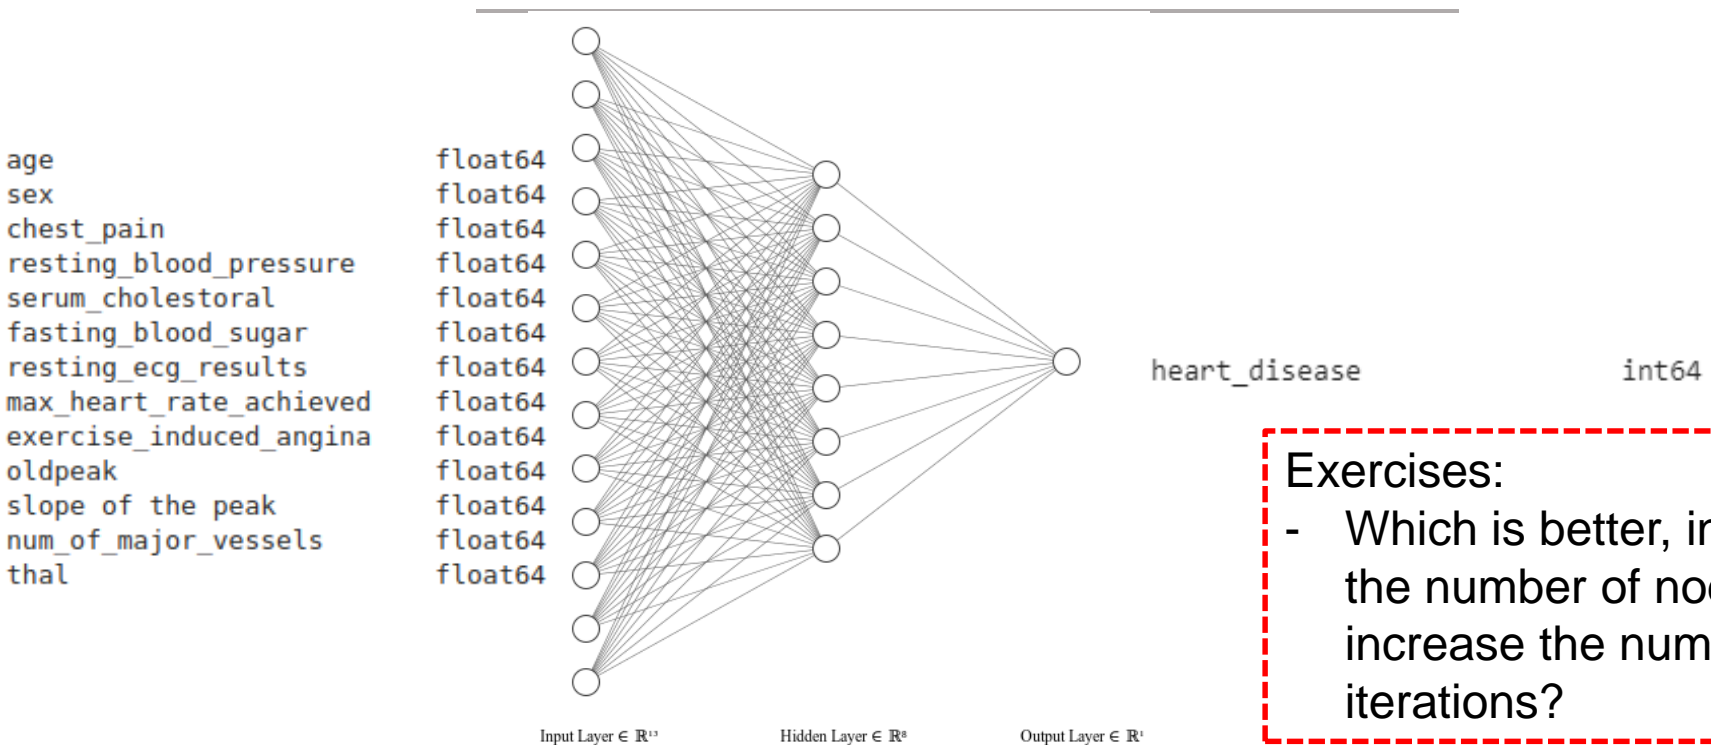
<https://www.youtube.com/watch?v=tIeHLnjs5U8>



***15 Mins  
Break***



# Activity 4 – Building NN with Python



## Exercises:

- Which is better, increase the number of nodes or increase the number of iterations?

### Step 1:

Watch and listen to the instructor's demonstration



15 mins

### Step 2:

Work through the activities



20 mins

**Individual Activity**



# Keras

---

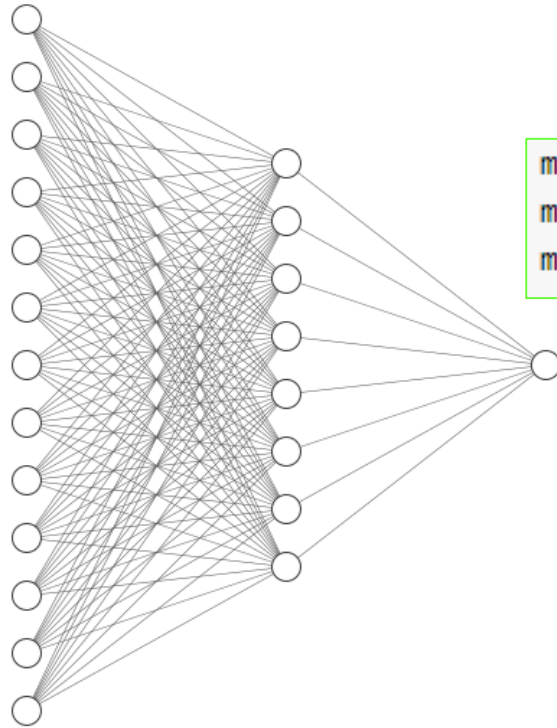
- Keras is one of the 5 most popular deep learning framework
- Built on Tensorflow 2.0 which allows Keras to run on GPUs as well as TPUs.
- It is very easy to use and create deep neural networks
- Homepage: <https://keras.io/>
- API reference: <https://keras.io/api/>







# Activity 5 – Building NN with Keras



Input Layer  $\in \mathbb{R}^{13}$

Hidden Layer  $\in \mathbb{R}^8$

Output Layer  $\in \mathbb{R}^1$

```
model = Sequential()  
model.add(Dense(8, input_shape=(13,), activation='relu'))  
model.add(Dense(1, activation='sigmoid'))
```

## Exercises:

- Add a hidden layer with 10 nodes
- Change the optimizer

### Step 1:

Watch and listen to the instructor's demonstration



10 mins

### Step 2:

Work through the activities



15 mins

**Individual Activity**

# 60 mins Lunch Break

Lunch break xx:xx – yy:yy

**LUNCH BREAK**




# Image Convolution

---



# Image Convolution

- In image processing, convolution is the process of transforming an image by applying a kernel over each pixel and its local neighbors across the entire image. The kernel is a matrix of values whose size and values determine the transformation effect of the convolution process.





# Types of Kernel / Filters

- **Basic image pre-processing or enhancement**

- Sharpening, Brightness, Blurring, etc
- <https://medium.com/@bdhuma/6-basic-things-to-know-about-convolution-daef5e1bc411>

- **Feature extraction**

- Line detector
- edge detector


Vertical Line Detector

-1	2	-1
-1	2	-1
-1	2	-1

Horizontal Line Detector

-1	-1	-1
2	2	2
-1	-1	-1


Input image



Convolution  
Kernel

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Feature map







# Types of Kernel / Filters

- **Pattern extraction**

- Patterns, grids


Padded input



Padded kernel



Cross-correlated result



<https://developer.nvidia.com/discover/convolution>

- **Non-standard Kernels**


- Transpose, Dilated, Deformable
- <https://towardsdatascience.com/types-of-convolution-kernels-simplified-f040cb307c37>



# Grid Size

- The number of pixels a kernel “sees” at once
- Typically use odd numbers so that there is a “center” pixel
- Kernel does not need to be square


Height: 3, Width: 3



Height: 1, Width: 3



Height: 3, Width: 1





# Padding

- Using Kernels directly, there will be an “edge effect”
- Pixels near the edge will not be used as “center pixels” since there are not enough surrounding pixels
- Padding adds extra pixels around the frame
- So every pixel of the original image will be a center pixel as the kernel moves across the image

0	0	0	0	0	0	0	0
0	3	3	4	4	7	0	0
0	9	7	6	5	8	2	0
0	6	5	5	6	9	2	0
0	7	1	3	2	7	8	0
0	0	3	7	1	8	3	0
0	4	0	4	3	2	2	0
0	0	0	0	0	0	0	0

\*

1	0	-1
1	0	-1
1	0	-1

 $3 \times 3$ 

=

-10	-13	1			
-9	3	0			

 $6 \times 6$  $6 \times 6 \rightarrow 8 \times 8$

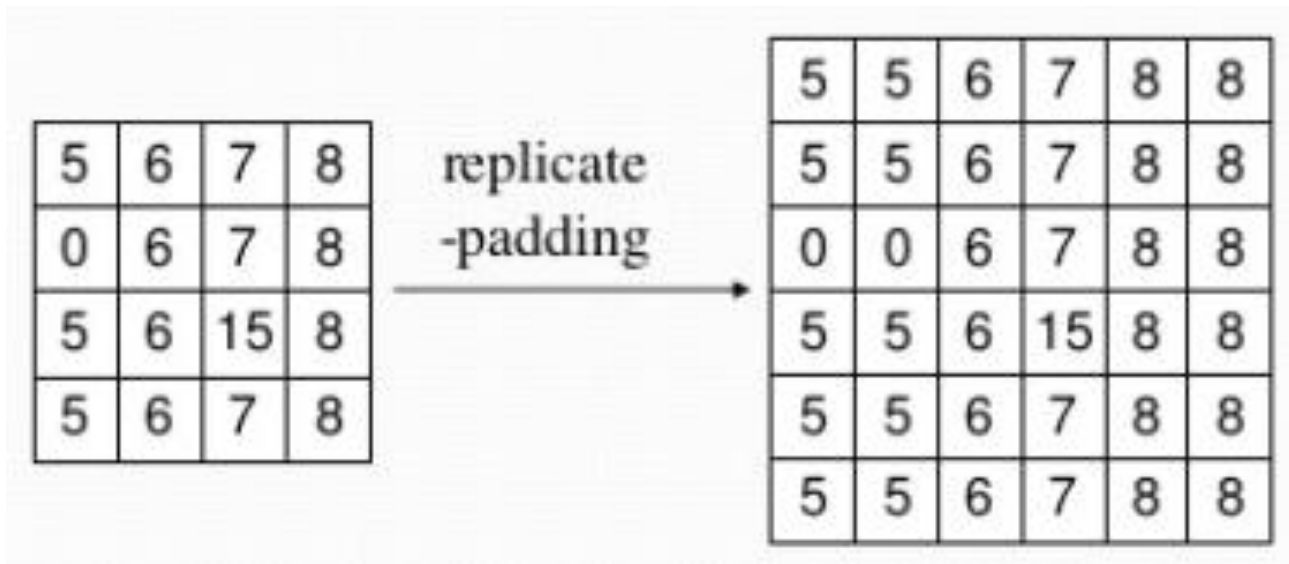




# Padding

- **Three form of padding**

- Zero padding – Added pixels with a value of 0
- One padding – Added pixels with a value of 1
- Replicate padding – Added pixels with the edge values of the image



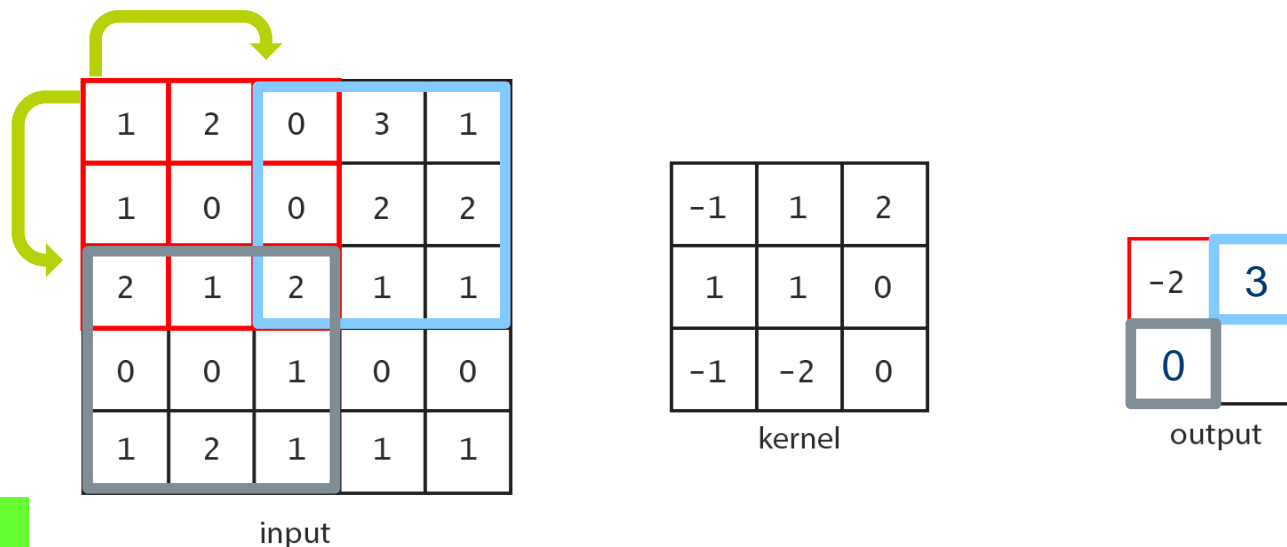
<https://pt.slideshare.net/zertux/nonlinear-filtering>



# Stride

- The "step size" as the kernel moves across the image
- Can be different for vertical and horizontal steps (but usually is the same value)
- When stride is greater than 1, it scales down the output dimension

## Stride 2 Example – No Padding






# Depth

---

- **In images, each pixel may be represented by multiple values. This is also known as image mode.**
- **The number of values is referred to as “channels”**
  - RGB image – 3 channels
  - CMYK – 4 channels
- **Kernel of the same “depth” in order to convolute the input channels.**
- **The output of the convolution will have the same depth or channel.**

# More on Image Convolution (optional)

---



<https://www.youtube.com/watch?v=8rrHTtUzyZA>

# Convolution Neural Network

---




# Basic Idea

- The image convolution formula is very similar to the output formula of neural network.

Image Convolution	Maps to	Neural Network
Input	→	Input nodes
Kernel	→	Weights
Output	→	Output nodes

$$\begin{array}{|c|c|c|} \hline \text{Input} & & \\ \hline 3 & 2 & 1 \\ \hline 1 & 2 & 3 \\ \hline 1 & 1 & 1 \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline \text{Kernel} & & \\ \hline -1 & 0 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & -2 & \\ \hline & 2 & \\ \hline & 0 & \\ \hline \end{array}$$





# Basic Idea

---


- The network is not fully connected. Each output node (pixel) is only connected to its convolve input nodes (pixels)
- Same set of weights (kernels) across the entire image
- Training of the weights can be applied
  - The trained weights form the kernel that is best to analyse the images
- A **convolution layer** consists of n-number of kernels. Each kernel will convolve with the input image(s) to produce n-number of output images.
- The activation function of a convolution layer is Relu because (i) there is no requirement of 'thresholding' the convoluted output image, and (ii) there is no negative values in an image.



# Other types of layers

- **Pooling layer**

- Shrinks the dimensions of an image (or reduce its size) by mapping a patch of pixels to one value.
  - Max-pooling
  - Average-pooling



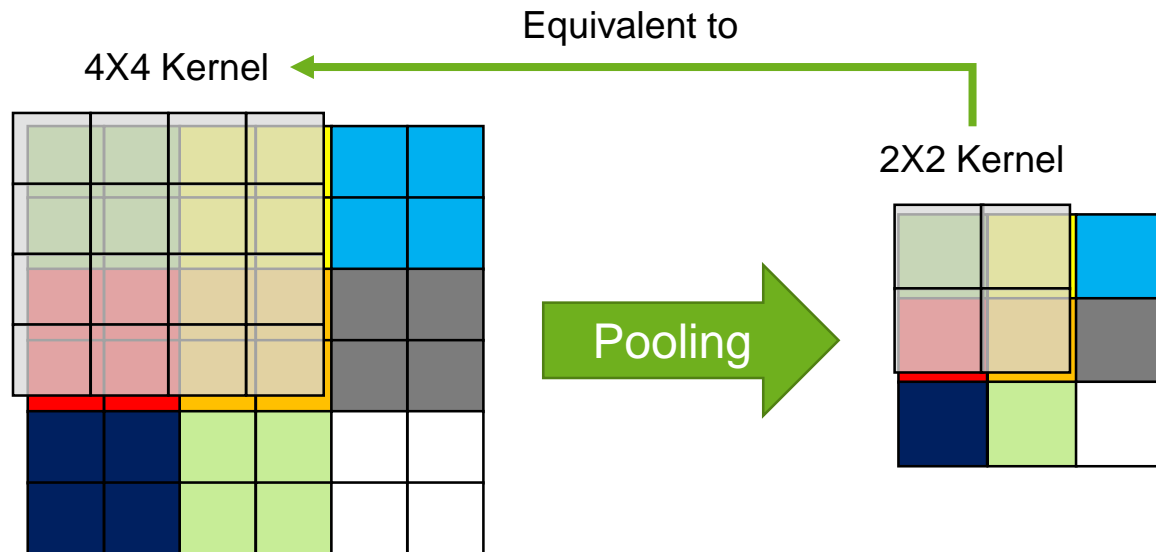




# Other types of layers

- **Pooling layer**

- To enable scale-invariant and shift-invariant
- Simulate convolution of bigger dimension







# Other types of layers

- **Dropout layer**

- Randomly removes one or more nodes from a network
- To prevent overfitting of the model
- Dropout is **ONLY** performed during training



(a) Standard Neural Net



(b) After applying dropout.

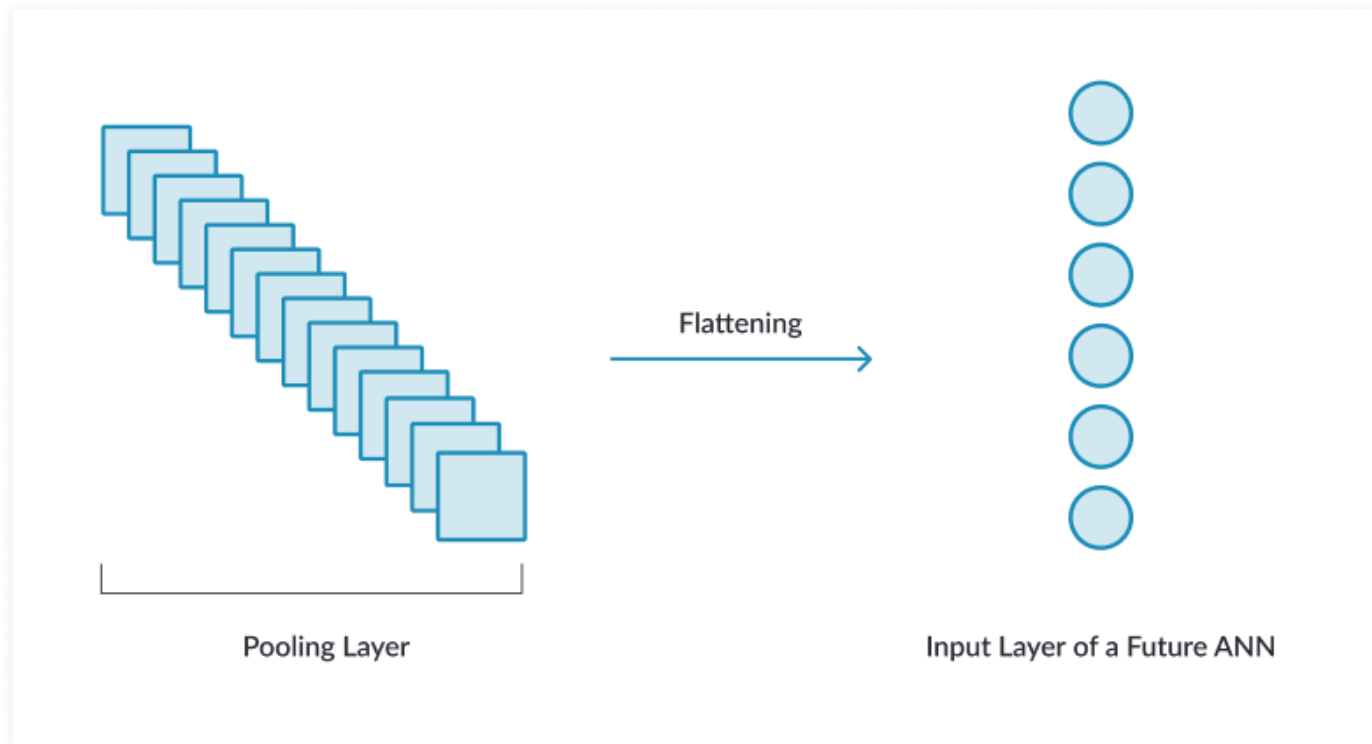
Ref: <https://medium.com/analytics-vidhya/a-simple-introduction-to-dropout-regularization-with-code-5279489dda1e>



# Other types of layers

- **Flatten layer**

- transforms a two-dimensional matrix of features into a vector that can be fed into a fully connected neural network classifier

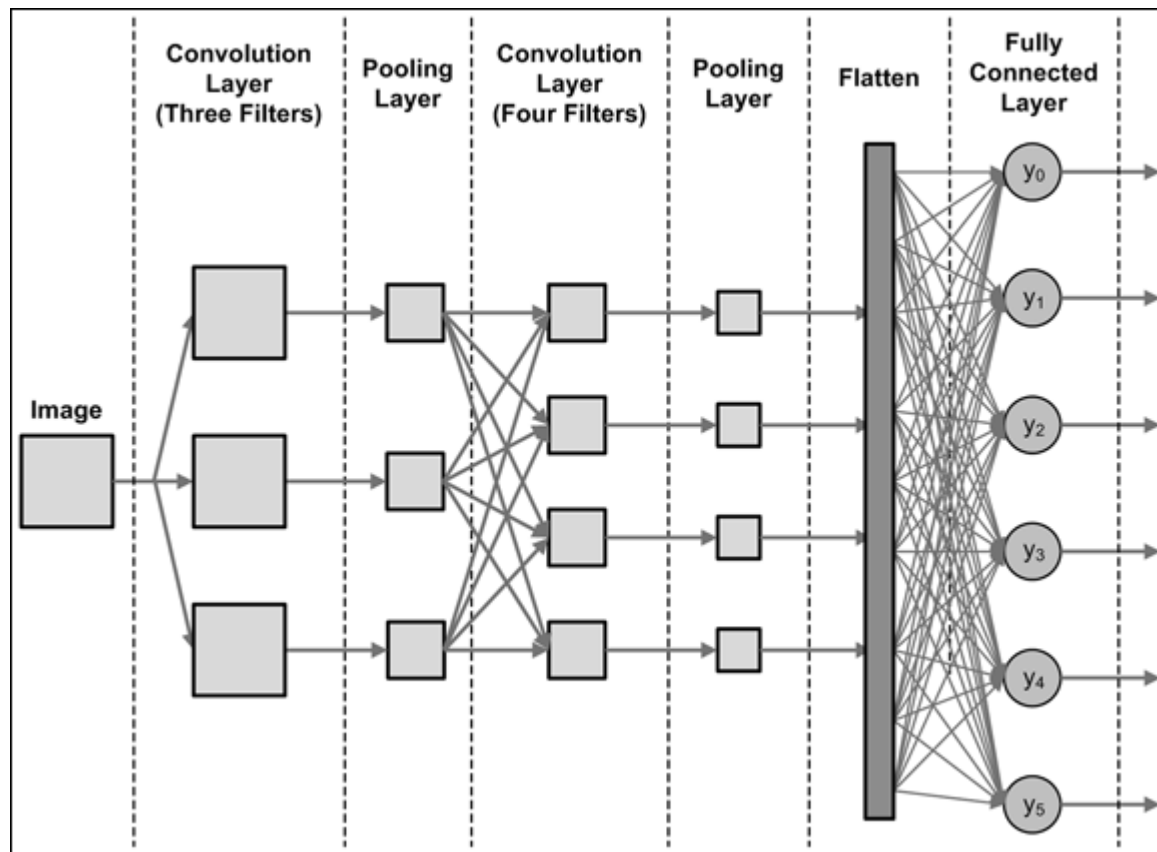


<https://missinglink.ai/guides/keras/using-keras-flatten-operation-cnn-models-code-examples/>



# Forming the CNN

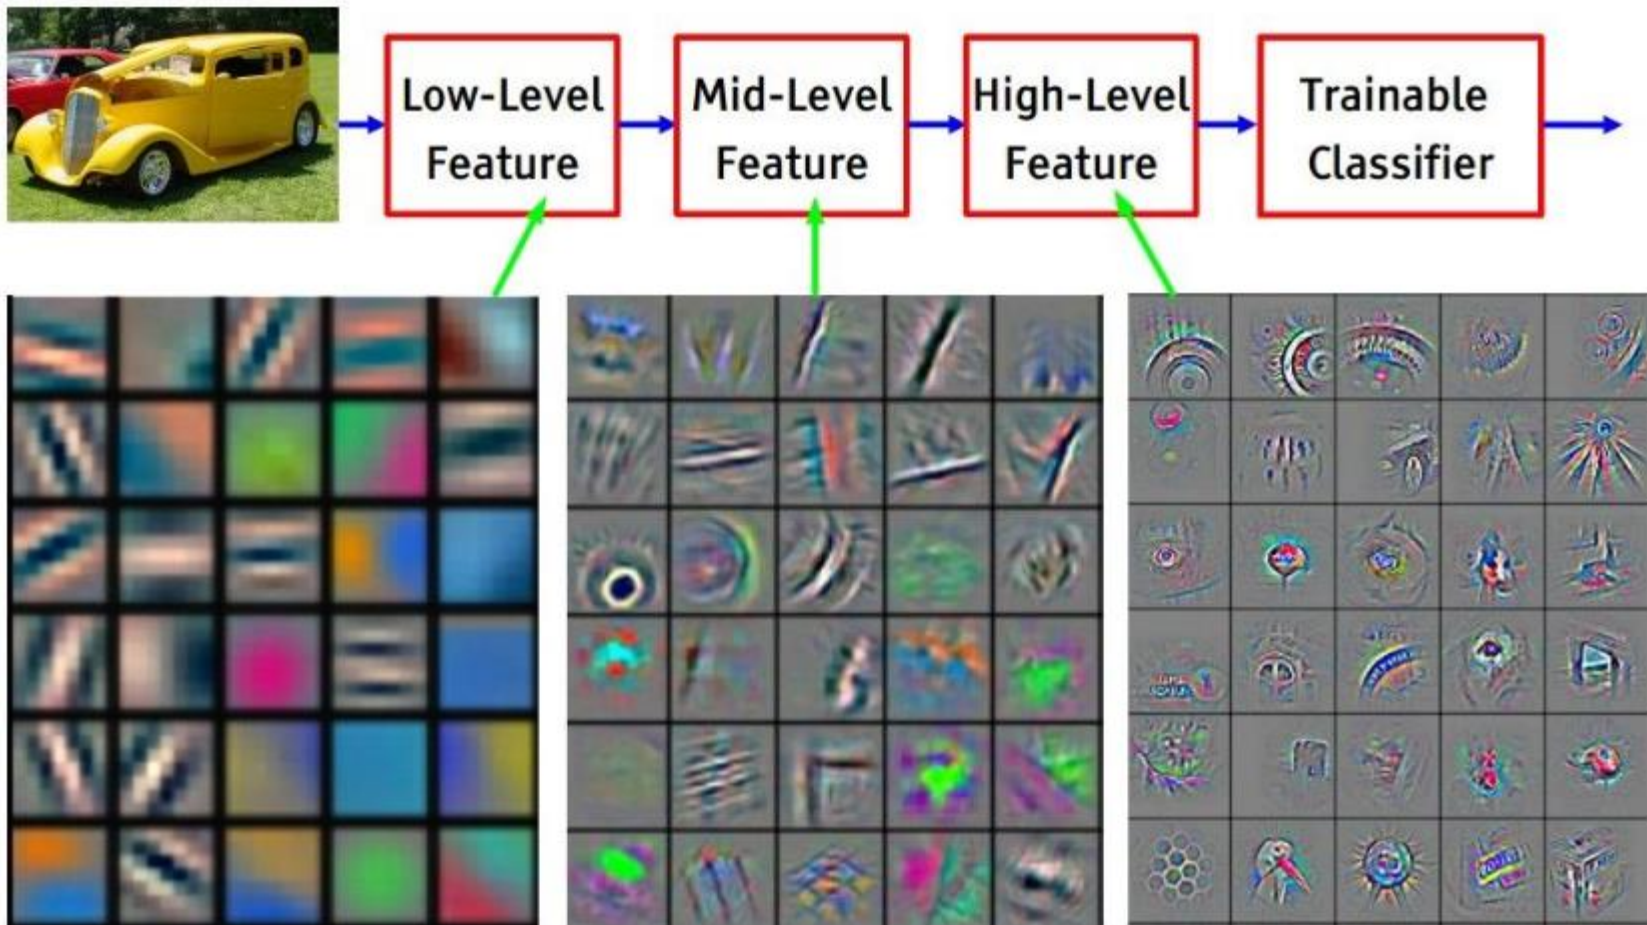
- By combining the convolution, pooling, flatten and dense (fully connected) layers, they formed a Convolution Neural Network.





# How CNN works

- Each layer extracts the **features** with increasing complexity





# How CNN works

- The dense layers (classifier) puts these features together to determine what is the image.




Ref: <https://towardsdatascience.com/simple-introduction-to-convolutional-neural-networks-cdf8d3077bac>



# How CNN works (Optional)

---

- **Explain by YouTube channel – deeplizard**
  - Check out the series on CNN
  - [https://www.youtube.com/playlist?list=PLZbbT5o\\_s2xq7Lwl2y8\\_QtvuXZedL6tQU](https://www.youtube.com/playlist?list=PLZbbT5o_s2xq7Lwl2y8_QtvuXZedL6tQU)



[https://www.youtube.com/watch?v=YRhxdVk\\_sls](https://www.youtube.com/watch?v=YRhxdVk_sls)





***15 Mins  
Break***







# Survey

---

- **Scan this QR-Code for the course feedback survey**



# Quiz

---

- **Go to this link for the quiz:**